

Lab-6

Iris Classify

1. Import libraries:

```
In [79]: #import all library
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

2. Loading the dataset.:

Load dataset

```
In [80]: iris_dataset = pd.read_csv('D:\Iris ML\dataset\Iris1.csv')
```

3. Summarizing the dataset:

- Dimensions of the dataset.

dimension

```
In [81]: iris_dataset.shape
```

```
Out[81]: (150, 6)
```

- Peek at the data itself.

peek

```
In [82]: iris_dataset.head()
```

```
Out[82]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

- Statistical summary of all attributes.

Statistical summary

```
In [84]: iris_dataset.describe()
```

```
Out[84]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

- Breakdown of the data by the class variable.

Breakdown by class variable

```
In [85]: iris_dataset['Id'].value_counts()
```

```
Out[85]: Id
1         1
95        1
97        1
98        1
99        1
..
51        1
52        1
53        1
54        1
150       1
Name: count, Length: 150, dtype: int64
```

```
In [86]: iris_dataset['SepalLengthCm'].value_counts()
```

```
Out[86]: SepalLengthCm
5.0      10
5.1       9
6.3       9
5.7       8
6.7       8
5.8       7
5.5       7
6.4       7
4.9       6
5.4       6
6.1       6
6.0       6
5.6       6
4.8       5
6.5       5
6.2       4
7.7       4
6.9       4
4.6       4
5.2       4
5.9       3
4.4       3
7.2       3
6.8       3
6.6       2
4.7       2
7.6       1
7.4       1
7.3       1
7.0       1
7.1       1
5.3       1
4.3       1
4.5       1
7.9       1
Name: count, dtype: int64
```

```
In [87]: iris_dataset['SepalWidthCm'].value_counts()
```

```
Out[87]: SepalWidthCm
3.0      26
2.8      14
3.2      13
3.1      12
3.4      12
2.9      10
2.7       9
2.5       8
3.5       6
3.3       6
3.8       6
2.6       5
2.3       4
3.7       3
2.4       3
2.2       3
3.6       3
3.9       2
4.4       1
4.0       1
4.1       1
4.2       1
2.0       1
Name: count, dtype: int64
```

```
In [88]: iris_dataset['PetalLengthCm'].value_counts()
```

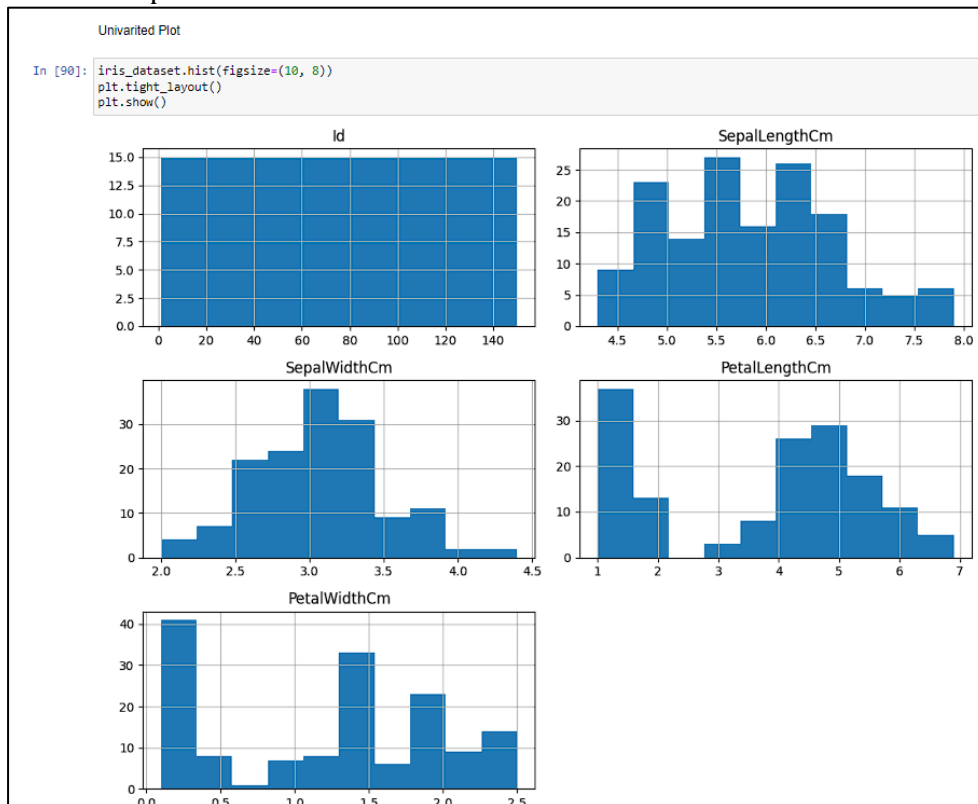
```
Out[88]: PetalLengthCm
1.5      14
1.4      12
5.1       8
4.5       8
1.6       7
1.3       7
5.6       6
4.7       5
4.9       5
4.0       5
4.2       4
5.0       4
4.4       4
4.8       4
1.7       4
3.9       3
4.6       3
5.7       3
4.1       3
5.5       3
6.1       3
5.8       3
3.3       2
5.4       2
6.7       2
5.3       2
5.9       2
6.0       2
1.2       2
4.3       2
1.9       2
3.5       2
5.2       2
3.0       1
1.1       1
3.7       1
3.8       1
6.6       1
6.3       1
1.0       1
6.9       1
3.6       1
6.4       1
Name: count, dtype: int64
```

```
In [89]: iris_dataset['PetalWidthCm'].value_counts()
```

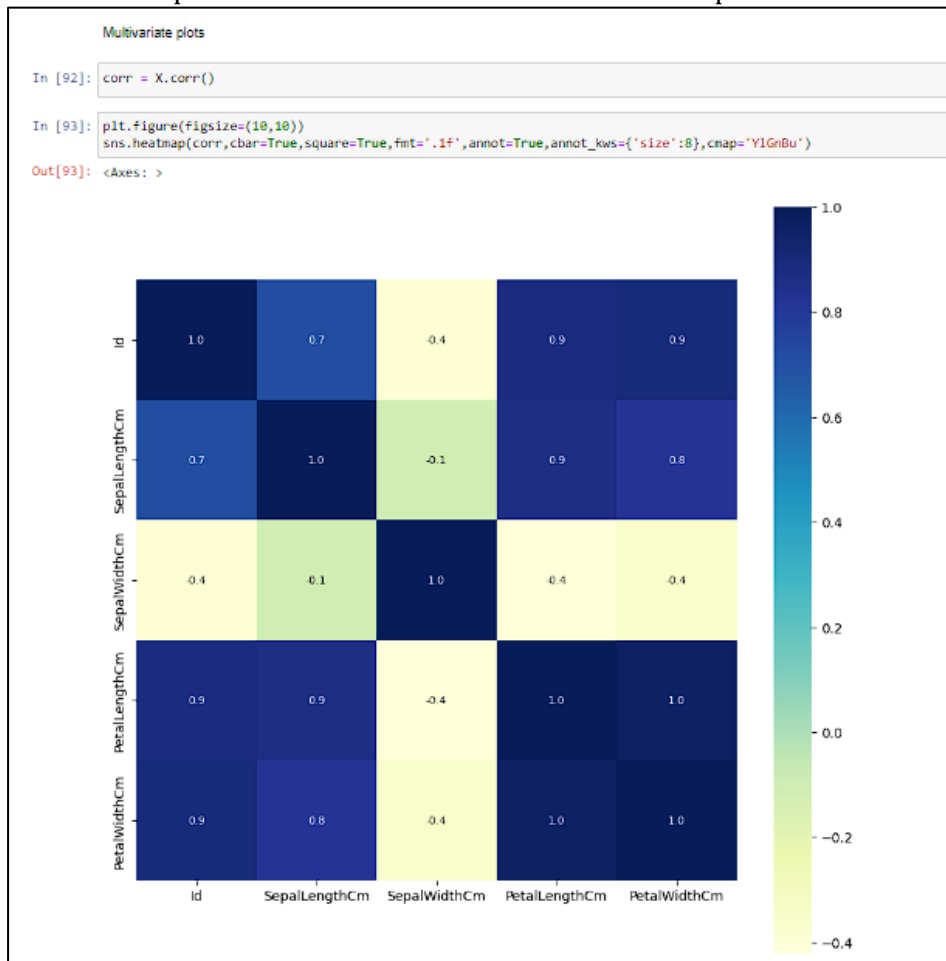
```
Out[89]: PetalWidthCm
0.2      28
1.3      13
1.8      12
1.5      12
1.4       8
2.3       8
1.0       7
0.4       7
0.3       7
0.1       6
2.1       6
2.0       6
1.2       5
1.9       5
1.6       4
2.5       3
2.2       3
2.4       3
1.1       3
1.7       2
0.6       1
0.5       1
Name: count, dtype: int64
```

4. Visualizing the dataset.

- Univariate plots to better understand each attribute.



- Multivariate plots to better understand the relationships between attributes.



5. Evaluating some algorithms.

- Separate out a validation dataset.

```
In [94]: y = iris_dataset['Species']

In [95]: print(y)

0      Iris-setosa
1      Iris-setosa
2      Iris-setosa
3      Iris-setosa
4      Iris-setosa
...
145     Iris-virginica
146     Iris-virginica
147     Iris-virginica
148     Iris-virginica
149     Iris-virginica
Name: Species, Length: 150, dtype: object
```

Split the data

```
In [96]: X_train, X_validation, y_train, y_validation = train_test_split(X, y, test_size=0.2, random_state=1, stratify=y)

In [97]: print(y)

0      Iris-setosa
1      Iris-setosa
2      Iris-setosa
3      Iris-setosa
4      Iris-setosa
...
145     Iris-virginica
146     Iris-virginica
147     Iris-virginica
148     Iris-virginica
149     Iris-virginica
Name: Species, Length: 150, dtype: object

In [98]: print(y.shape,y_train.shape,y_validation.shape)

(150,) (120,) (30,)

In [99]: print(X)

   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm
0    1             5.1           3.5           1.4           0.2
1    2             4.9           3.0           1.4           0.2
2    3             4.7           3.2           1.3           0.2
3    4             4.6           3.1           1.5           0.2
4    5             5.0           3.6           1.4           0.2
..  ...
145  146             6.7           3.0           5.2           2.3
146  147             6.3           2.5           5.0           1.9
147  148             6.5           3.0           5.2           2.0
148  149             6.2           3.4           5.4           2.3
149  150             5.9           3.0           5.1           1.8

[150 rows x 5 columns]

In [100]: print(X.shape,X_train.shape,X_validation.shape)
```

- Set-up the test harness to use 10-fold cross validation.

Set-up the test harness to use 10-fold cross validation.

```
In [101]: # Set-up 10-fold cross-validation
kfold = StratifiedKFold(n_splits=10, random_state=1, shuffle=True)
```

- Build multiple different models to predict species from flower measurements

1. Logistic Regression (LR)

Logistic Regression (LR)

```
In [102]: # Create a Logistic Regression model
model = LogisticRegression(solver='liblinear', multi_class='ovr')

In [103]: # Perform 10-fold cross-validation and evaluate the model
cv_results = cross_val_score(model, X_train, y_train, cv=kfold, scoring='accuracy')

# Output cross-validation results
print(f"Logistic Regression: {cv_results.mean():.4f} ({cv_results.std():.4f})")

Logistic Regression: 0.9333 (0.0624)
```

```
In [104]: # Train the Logistic Regression model on the training dataset
model.fit(X_train, y_train)
```

```
Out[104]: LogisticRegression(multi_class='ovr', solver='liblinear')
```

```
In [105]: # Make predictions on the validation dataset
predictions = model.predict(X_validation)

# Evaluate accuracy on the validation dataset
print(f"Accuracy on validation set: {accuracy_score(y_validation, predictions):.4f}")

Accuracy on validation set: 0.9000
```

2. Linear Discriminant Analysis (LDA)

Linear Discriminant Analysis

```
In [106]: lda = LinearDiscriminantAnalysis()

In [107]: # Perform 10-fold cross-validation on the training dataset
cv_results = cross_val_score(lda, X_train, y_train, cv=kfold, scoring='accuracy')

In [108]: # Print cross-validation results
print(f"LDA: {cv_results.mean():.4f} ({cv_results.std():.4f})")

LDA: 1.0000 (0.0000)

In [109]: # Train the LDA model on the entire training dataset
lda.fit(X_train, y_train)

Out[109]: LinearDiscriminantAnalysis()
```

```
In [110]: # Make predictions on the validation dataset
predictions = lda.predict(X_validation)

In [111]: # Evaluate the accuracy on the validation dataset
accuracy = accuracy_score(y_validation, predictions)
print(f"Accuracy on validation set: {accuracy:.4f}")

Accuracy on validation set: 1.0000
```

3. K-Nearest Neighbors (KNN).

```
In [112]: # Initialize the KNN model
knn = KNeighborsClassifier()

In [113]: # Evaluate KNN model using 10-fold cross-validation
cv_results = cross_val_score(knn, X_train, y_train, cv=kfold, scoring='accuracy')

# Print the cross-validation accuracy results
print(f"KNN Cross-Validation Accuracy: {cv_results.mean():.4f} ({cv_results.std():.4f})")

KNN Cross-Validation Accuracy: 1.0000 (0.0000)

In [114]: # Train the KNN model on the full training dataset
knn.fit(X_train, y_train)

Out[114]: KNeighborsClassifier()
```

```
In [115]: # Make predictions on the validation dataset
predictions = knn.predict(X_validation)

# Evaluate the model performance on the validation dataset
accuracy = accuracy_score(y_validation, predictions)
print(f"Accuracy on Validation Set: {accuracy:.4f}")

Accuracy on Validation Set: 1.0000
```

4. Classification and Regression Trees (CART).

```
In [120]: # Initialize the CART model
cart_model = DecisionTreeClassifier()

In [121]: # Display cross-validation results
print(f"CART - Cross-Validation Accuracy: {cv_results.mean():.4f} ({cv_results.std():.4f})")

CART - Cross-Validation Accuracy: 1.0000 (0.0000)

In [122]: # Train the CART model on the entire training dataset
cart_model.fit(X_train, y_train)

Out[122]: DecisionTreeClassifier()

In [123]: # Make predictions on the validation dataset
predictions = cart_model.predict(X_validation)

# Evaluate the performance on the validation dataset
validation_accuracy = accuracy_score(y_validation, predictions)
print(f"CART - Accuracy on validation set: {validation_accuracy:.4f}")

CART - Accuracy on validation set: 1.0000
```

5. Gaussian Naive Bayes (NB).

```
In [124]: # Initialize Gaussian Naive Bayes model
model = GaussianNB()

In [125]: print(f"Gaussian Naive Bayes cross-validation accuracy: {cv_results.mean():.4f} ({cv_results.std():.4f})")

Gaussian Naive Bayes cross-validation accuracy: 1.0000 (0.0000)

In [126]: # Fit the model on the training dataset
model.fit(X_train, y_train)

Out[126]: GaussianNB()

In [127]: # Make predictions on the validation dataset
predictions = model.predict(X_validation)

# Evaluate accuracy on the validation dataset
accuracy = accuracy_score(y_validation, predictions)
print(f"Accuracy on validation set: {accuracy:.4f}")

Accuracy on validation set: 0.9667
```

6. Support Vector Machines (SVM).

Support Vector Machines (SVM).

```
In [116]: # Initialize the SVM model
model = SVC(gamma='auto')

In [117]: print(f"SVM: {cv_results.mean():.4f} ({cv_results.std():.4f})")

SVM: 1.0000 (0.0000)

In [118]: # Train the SVM model on the training dataset
model.fit(X_train, y_train)

Out[118]: SVC(gamma='auto')
```

```
In [119]: # Make predictions on the validation dataset
          predictions = model.predict(X_validation)

          # Evaluate the accuracy on the validation dataset
          print(f"Accuracy on validation set: {accuracy_score(y_validation, predictions):.4f}")

          Accuracy on validation set: 1.0000
```

- Accuracy of the models:
 - Logistic Regression (LR): **0.9000**
 - Linear Discriminant Analysis (LDA):**1.000**
 - K-Nearest Neighbors (KNN).**1.000**
 - Classification and Regression Trees (CART).: **1.000**
 - Gaussian Naive Bayes (NB).**0.9667**
 - Support Vector Machines (SVM).: **1.000**